

B2_Statistics_Modeling_and_Inference

June 14, 2026

1 Credit Risk Statistics, Modeling, and Inference

Notebook	B2_Statistics_Modeling_and_Inference.ipynb
Author	N L N Sai Krishna Akula
Course	AAI 500 - Applied Statistics & AI
Dataset	Credit Risk Modeling Dataset
Date	June 3, 2026
GitHub	https://github.com/diaeshantony-usd/applied-statistics-and-ai-team-project

This notebook starts from the saved Phase 2 modeling dataset. It runs inferential statistics, Bayesian Network modeling, ML classification, ML regression, model saving, and final applicant inference.

```
[47]: # Core data libraries
import warnings
import joblib
from pathlib import Path

import numpy as np
import pandas as pd

# Visualization libraries
import matplotlib.pyplot as plt
import networkx as nx
import seaborn as sns

# Inferential statistics
from scipy import stats

# Scikit-learn preprocessing
from sklearn.base import clone
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, StandardScaler
```

```

# Scikit-learn model selection and metrics
from sklearn.model_selection import StratifiedKFold, cross_val_predict, \
    cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import (
    accuracy_score,
    classification_report,
    f1_score,
    mean_absolute_error,
    mean_squared_error,
    precision_score,
    r2_score,
    recall_score,
    roc_auc_score,
)

# Scikit-learn estimators
from sklearn.ensemble import GradientBoostingClassifier, \
    GradientBoostingRegressor
from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
from sklearn.linear_model import LinearRegression, LogisticRegression, Ridge
from catboost import CatBoostClassifier
from catboost import CatBoostRegressor

warnings.filterwarnings("ignore")
sns.set_theme(style="whitegrid")
pd.set_option("display.max_columns", 100)

RANDOM_STATE = 42
TEST_SIZE = 0.20
PHASE2_DATA_PATH = "../data/processed/phase2/credit_risk_cleaned.csv"
CLASSIFICATION_MODEL_DIR = Path("../data/models/classification")
REGRESSION_MODEL_DIR = Path("../data/models/regression")
BAYESIAN_MODEL_DIR = Path("../data/models/bayesian_network")

print("Libraries loaded successfully.")
print("Phase 2 dataset configured:", PHASE2_DATA_PATH)
print("Classification model folder:", CLASSIFICATION_MODEL_DIR)
print("Regression model folder:", REGRESSION_MODEL_DIR)
print("Bayesian Network model folder:", BAYESIAN_MODEL_DIR)

```

```

Libraries loaded successfully.
Phase 2 dataset configured: ../data/processed/phase2/credit_risk_cleaned.csv
Classification model folder: ../data\models\classification
Regression model folder: ../data\models\regression
Bayesian Network model folder: ../data\models\bayesian_network

```

1.1 Setup for Statistics and Modeling

This notebook starts from the saved Phase 2 CSV created in Notebook 1. The setup cell includes the imports, paths, and helper functions needed to run the statistics, modeling, saved-model loading, and applicant inference steps.

```
[48]: def add_credit_features(data):
    # Copy the input so the original cleaned dataset is not modified directly.
    out = data.copy()

    # Log transforms reduce the influence of very large income and loan values.
    out["log_person_income"] = np.log1p(out["person_income"])
    out["log_loan_amnt"] = np.log1p(out["loan_amnt"])

    # Ratio features describe affordability and borrower history in relative
    # terms.
    out["income_to_loan_ratio"] = out["person_income"] / out["loan_amnt"].
    clip(lower=1)
    out["employment_to_age_ratio"] = out["person_emp_length"] /
    out["person_age"].clip(lower=1)
    out["credit_history_to_age_ratio"] = out["cb_person_cred_hist_length"] /
    out["person_age"].clip(lower=1)
    out["default_on_file_binary"] = (out["cb_person_default_on_file"] == "Y").
    astype(int)

    # Numeric grade is useful for analysis while model preprocessing still
    # handles grade encoding.
    grade_map = {"A": 0, "B": 1, "C": 2, "D": 3, "E": 4, "F": 5, "G": 6}
    out["loan_grade_numeric"] = out["loan_grade"].map(grade_map)

    # Grouped features make borrower and loan risk patterns easier to explain.
    out["income_group"] = pd.cut(
        out["person_income"],
        bins=[0, 25000, 50000, 75000, 100000, np.inf],
        labels=["Low", "Lower-Middle", "Middle", "Upper-Middle", "High"],
        include_lowest=True,
    )
    out["loan_amount_group"] = pd.cut(
        out["loan_amnt"],
        bins=[0, 5000, 10000, 20000, np.inf],
        labels=["Small", "Medium", "Large", "Very Large"],
        include_lowest=True,
    )
    out["loan_percent_income_group"] = pd.cut(
        out["loan_percent_income"],
        bins=[0, 0.10, 0.20, 0.35, np.inf],
```

```

        labels=["Low Burden", "Moderate Burden", "High Burden", "Very High_Burden"],
        include_lowest=True,
    )
    out["age_group"] = pd.cut(
        out["person_age"],
        bins=[0, 25, 35, np.inf],
        labels=["Young", "Adult", "Experienced Adult"],
        include_lowest=True,
    )
    # Return the feature-engineered dataframe for all later modeling steps.
    return out

```

1.2 3. Statistical and Bayesian Analysis

1.2.1 3.1 Inferential Statistics

Chapter 3 starts by loading the saved Phase 2 dataset. This makes the statistics and modeling steps reproducible without depending on earlier notebook cells already being in memory. Inferential statistics help check whether group differences look real or could just be random. This dataset supports mean-comparison tests for borrower, loan, and credit-history variables.

Welch's t-test is used for two-group comparisons between default and non-default loans because it does not assume equal variance between the groups. One-way ANOVA is used to compare numeric means across more than two credit grades.

H0 means the null hypothesis. It represents the assumption that there is no real difference between the groups being compared. If the p-value is below 0.05, the result rejects H0, meaning the observed difference is unlikely to be explained by random variation alone.

```

[49]: # Reload the saved modeling dataset before statistics and modeling.
df_model = pd.read_csv(PHASE2_DATA_PATH)
print("Modeling dataset loaded successfully.")
print("Shape:", df_model.shape)
df_model.head()

```

Modeling dataset loaded successfully.

Shape: (29567, 23)

```

[49]:   person_age  person_income  person_home_ownership  person_emp_length \
0          21          9600             OWN             5.0
1          25          9600             MORTGAGE           1.0
2          23         65500             RENT             4.0
3          24         54400             RENT             8.0
4          21          9900             OWN             2.0

   loan_intent  loan_grade  loan_amnt  loan_int_rate  loan_status \
0  EDUCATION          B      1000      11.14           0

```

1	MEDICAL	C	5500	12.87	1
2	MEDICAL	C	35000	15.23	1
3	MEDICAL	C	35000	14.27	1
4	VENTURE	A	2500	7.14	1

	loan_percent_income	cb_person_default_on_file	cb_person_cred_hist_length	\
0	0.10	N		2
1	0.57	N		3
2	0.53	N		2
3	0.55	Y		4
4	0.25	N		2

	log_person_income	log_loan_amnt	income_to_loan_ratio	\
0	9.169623	6.908755	9.600000	
1	9.169623	8.612685	1.745455	
2	11.089821	10.463132	1.871429	
3	10.904138	10.463132	1.554286	
4	9.200391	7.824446	3.960000	

	employment_to_age_ratio	credit_history_to_age_ratio	\
0	0.238095	0.095238	
1	0.040000	0.120000	
2	0.173913	0.086957	
3	0.333333	0.166667	
4	0.095238	0.095238	

	default_on_file_binary	loan_grade_numeric	income_group	loan_amount_group	\
0	0	1	Low	Small	
1	0	2	Low	Medium	
2	0	2	Middle	Very Large	
3	1	2	Middle	Very Large	
4	0	0	Low	Small	

	loan_percent_income_group	age_group
0	Low Burden	Young
1	Very High Burden	Young
2	Very High Burden	Young
3	Very High Burden	Young
4	High Burden	Young

[50]: ALPHA = 0.05

```
ttest_numeric_features = [
    "person_income", "loan_amnt", "loan_percent_income",
    "loan_int_rate", "person_emp_length", "cb_person_cred_hist_length",
]
```

```

def run_welch_tests(data, features, target_col="loan_status", alpha=ALPHA):
    # Welch's t-test compares numeric means between default and non-default
    ↪loans.
    rows = []
    for feature in features:
        non_default_values = data.loc[data[target_col] == 0, feature].dropna()
        default_values = data.loc[data[target_col] == 1, feature].dropna()

        # equal_var=False applies Welch's version of the independent t-test.
        statistic, p_value = stats.ttest_ind(
            non_default_values,
            default_values,
            equal_var=False,
            nan_policy="omit",
        )
        decision = "Reject H0" if p_value < alpha else "Fail to reject H0"
        interpretation = (
            "Reject H0: evidence suggests the mean differs by default status."
            if p_value < alpha
            else "Fail to reject H0: no statistically significant mean
            ↪difference found."
        )
        rows.append({
            "test": "Welch t-test",
            "feature": feature,
            "h0": f"The mean {feature} is the same for default and non-default
            ↪loans.",
            "statistic": statistic,
            "p_value": p_value,
            "decision": decision,
            "interpretation": interpretation,
        })
    return pd.DataFrame(rows)

welch_results = run_welch_tests(df_model, ttest_numeric_features)

# Show Welch test results for default vs non-default groups.
welch_results.style.format({"statistic": "{:.4f}", "p_value": "{:.6f}"})

```

[50]: <pandas.io.formats.style.Styler at 0x1a024b25460>

For Welch's t-tests, H0 means the average numeric value is the same for default and non-default loans. When the p-value is below 0.05, H0 is rejected, meaning the feature shows a statistically meaningful difference between the two credit-outcome groups.

```

[51]: anova_numeric_features = [
        "person_income", "loan_amnt", "loan_percent_income",
        "loan_int_rate", "person_emp_length", "cb_person_cred_hist_length",
    ]

def run_grade_anova(data, features, group_col="loan_grade", alpha=ALPHA):
    # One-way ANOVA compares numeric means across all loan-grade groups.
    rows = []
    for feature in features:
        grouped_values = [
            group[feature].dropna()
            for _, group in data.groupby(group_col)
            if group[feature].dropna().shape[0] > 1
        ]

        # ANOVA requires at least two groups with enough observations.
        if len(grouped_values) >= 2:
            statistic, p_value = stats.f_oneway(*grouped_values)
        else:
            statistic, p_value = np.nan, np.nan

        decision = "Reject H0" if pd.notna(p_value) and p_value < alpha else
        ↪ "Fail to reject H0"
        interpretation = (
            "Reject H0: at least one grade group has a different mean."
            if pd.notna(p_value) and p_value < alpha
            else "Fail to reject H0: no statistically significant grade-level
            ↪ mean difference found."
        )
        rows.append({
            "test": "One-way ANOVA",
            "feature": feature,
            "h0": f"The mean {feature} is the same across all loan grades.",
            "statistic": statistic,
            "p_value": p_value,
            "decision": decision,
            "interpretation": interpretation,
        })
    return pd.DataFrame(rows)

anova_results = run_grade_anova(df_model, anova_numeric_features)

# Show ANOVA results for numeric variables across loan grades.
anova_results.style.format({"statistic": "{:.4f}", "p_value": "{:.6f}"})

```

[51]: <pandas.io.formats.style.Styler at 0x1a047f1aab0>

For ANOVA, H_0 means the average numeric value is the same across all loan-grade groups. A rejected H_0 means at least one grade group has a different average, which supports using grade-related patterns for model selection and interpretation.

1.2.2 3.2 Bayesian Network Model

A Bayesian Network shows how grouped variables are connected using probabilities. For this dataset, the network can be learned from grouped borrower and loan variables such as income group, loan burden group, prior default flag, loan intent, and loan grade.

This section is used for probabilistic structure interpretation. Because the data is observational, the learned DAG should not be treated as proof of real-world causality.

The fitted Bayesian Network can also estimate default probability for test records. It is used as a supporting model, while the Scikit-learn models remain the main prediction models because they use more detailed numeric information.

The fitted Bayesian Network is saved with metadata so it can be loaded again for applicant-level prediction.

```
[52]: # Run this focused analysis step and display the result for project review.
from pgmpy.inference import VariableElimination
from pgmpy.estimators import HillClimbSearch, BIC
from pgmpy.models import DiscreteBayesianNetwork as PgmpyBayesianNetwork
print("pgmpy imported successfully. Bayesian Network structure learning is_
↳available.")
```

pgmpy imported successfully. Bayesian Network structure learning is available.

```
[53]: bn_cols = [
    "income_group", "loan_amount_group", "loan_percent_income_group",
    "age_group", "person_home_ownership", "loan_intent",
    "cb_person_default_on_file", "loan_grade", "loan_status",
]

# Use only discrete/grouped variables because Bayesian Networks in this section_
↳are categorical.
bn_data = df_model[bn_cols].copy()
bn_data["loan_status"] = bn_data["loan_status"].map({0: "Non-Default", 1:
↳"Default"})
bn_data = bn_data.astype(str)

# Split the discrete Bayesian Network data for training and prediction_
↳evaluation.
bn_train_data, bn_test_data = train_test_split(
    bn_data,
    test_size=TEST_SIZE,
    stratify=bn_data["loan_status"],
```



```

    random_state=RANDOM_STATE,
)

# Use a training sample for structure learning so the section remains practical
↳to run.
#bn_learning_data = bn_train_data.sample(n=min(5000, len(bn_train_data)),
↳random_state=RANDOM_STATE)
bn_learning_data = bn_train_data
print("Bayesian Network training rows:", bn_train_data.shape[0])
print("Bayesian Network test rows:", bn_test_data.shape[0])
print("Bayesian Network structure-learning rows:", bn_learning_data.shape[0])
bn_learning_data.head()

```

Bayesian Network training rows: 23653

Bayesian Network test rows: 5914

Bayesian Network structure-learning rows: 23653

```

[53]:      income_group loan_amount_group loan_percent_income_group age_group \
19556          Low          Small          High Burden      Adult
17624          High          Large      Moderate Burden      Adult
14146          High          Large      Moderate Burden      Young
17846  Lower-Middle          Large    Very High Burden      Adult
2255          Low          Small          High Burden      Young

      person_home_ownership loan_intent cb_person_default_on_file loan_grade \
19556          RENT      VENTURE          Y          C
17624          RENT      MEDICAL          Y          F
14146      MORTGAGE      VENTURE          N          B
17846          RENT      PERSONAL          Y          C
2255          RENT      PERSONAL          Y          D

      loan_status
19556      Default
17624      Default
14146  Non-Default
17846      Default
2255      Default

```

```

[54]: bn_model = None
bn_edges = []
bn_metrics = None

# HillClimbSearch learns a DAG structure using a discrete scoring method.
hill_search = HillClimbSearch(bn_learning_data)
learned_structure = hill_search.estimate(
    scoring_method=BIC(bn_learning_data),

```

```

    max_indegree=3,
)

# Fit conditional probability tables on the learned graph.
bn_edges = list(learned_structure.edges())
bn_model = PgmpyBayesianNetwork(bn_edges)
bn_model.add_nodes_from(bn_cols)

bn_model.fit(bn_learning_data)
print("Bayesian Network parameters fitted with the default pgmpy discrete_
↪estimator.")

bn_edge_table = pd.DataFrame(bn_edges, columns=["from", "to"])
if bn_edge_table.empty:
    print("No edges were learned from the selected variables.")

bn_edge_table

```

```
0%|          | 14/1000000 [00:00<6:48:12, 40.83it/s]
```

Bayesian Network parameters fitted with the default pgmpy discrete estimator.

```

[54]:
      from      to
0      income_group  loan_percent_income_group
1      income_group      loan_amount_group
2      loan_amount_group  loan_percent_income_group
3      loan_amount_group      loan_status
4  loan_percent_income_group      loan_status
5      person_home_ownership  income_group
6      person_home_ownership      loan_status
7      person_home_ownership      loan_intent
8      person_home_ownership      loan_grade
9      loan_intent      age_group
10     loan_grade  cb_person_default_on_file
11     loan_status      loan_grade
12     loan_status      loan_intent
13     loan_status      age_group

```

The fitted Bayesian Network is saved after training. The metadata file stores the column list, target column, and fallback default probability needed for later prediction.

```

[55]: # Create the Bayesian Network model folder before saving files.
BAYESIAN_MODEL_DIR.mkdir(parents=True, exist_ok=True)

# Save the fitted Bayesian Network model.
bayesian_model_path = BAYESIAN_MODEL_DIR / "bayesian_network_model.joblib"

```

```

joblib.dump(bn_model, bayesian_model_path)

# Save metadata needed when the model is loaded later for prediction.
fallback_probability = float((bn_train_data["loan_status"] == "Default").mean())
bayesian_metadata_path = BAYESIAN_MODEL_DIR / "bayesian_network_metadata.joblib"
joblib.dump({
    "bn_cols": bn_cols,
    "fallback_probability": fallback_probability,
    "target_col": "loan_status",
}, bayesian_metadata_path)

bayesian_model_files = pd.DataFrame([
    {"artifact": "Bayesian Network model", "saved_path": ↳
    str(bayesian_model_path)},
    {"artifact": "Bayesian Network metadata", "saved_path": ↳
    str(bayesian_metadata_path)},
])
bayesian_model_files

```

```

[55]:          artifact \
0      Bayesian Network model
1  Bayesian Network metadata

          saved_path
0  ..\data\models\bayesian_network\bayesian_netwo...
1  ..\data\models\bayesian_network\bayesian_netwo...

```

The learned edge table lists probabilistic dependencies selected by HillClimbSearch. These edges describe relationships found in observational data and should be interpreted as dependency structure, not guaranteed causality.

```

[56]: # Plot the learned DAG so dependencies among grouped variables are visible.
dag = nx.DiGraph()
dag.add_nodes_from(bn_cols)
dag.add_edges_from(bn_edges)

plt.figure(figsize=(14, 6))
pos = nx.spring_layout(dag, seed=RANDOM_STATE, k=1.1)
nx.draw_networkx_nodes(
    dag,
    pos,
    node_color="lightsteelblue",
    edgecolors="navy",
    linewidths=1.2,
    node_size=1700,
)
nx.draw_networkx_edges(


```

```

dag,
pos,
arrows=True,
arrowstyle="->",
arrowsize=28,
edge_color="dimgray",
width=1.8,
min_source_margin=18,
min_target_margin=22,
connectionstyle="arc3,rad=0.08",
)
nx.draw_networkx_labels(dag, pos, font_size=8, font_weight="bold")
plt.title("Learned Bayesian Network DAG")
plt.axis("off")
plt.tight_layout()
plt.show()

```

Learned Bayesian Network DAG



```

[57]: def predict_default_probability_with_bn(model, row, target_col="loan_status",
    ↪ fallback_probability=None):
    # VariableElimination returns a probability table for loan_status using
    ↪ applicant evidence.
    inference = VariableElimination(model)
    if fallback_probability is None:
        fallback_probability = float((bn_train_data[target_col] == "Default").
    ↪ mean())
    evidence = {
        col: str(row[col])
        for col in bn_cols
        if col != target_col and col in model.nodes()
    }

```

```

try:
    query_result = inference.query(
        variables=[target_col],
        evidence=evidence,
        show_progress=False,
    )
    target_states = list(query_result.state_names[target_col])
    default_index = target_states.index("Default")
    predicted_probability = float(query_result.values[default_index])
    if not np.isfinite(predicted_probability):
        return fallback_probability
    return predicted_probability
except Exception:
    # If a rare evidence combination is unavailable, fall back to the
    ↪ training default rate.
    return fallback_probability

fallback_probability = float((bn_train_data["loan_status"] == "Default").mean())

# Predict default probability for each Bayesian Network test record.
bn_default_prob = bn_test_data.apply(
    lambda row: predict_default_probability_with_bn(bn_model, row),
    axis=1,
)

bn_default_prob = pd.to_numeric(bn_default_prob, errors="coerce")
bn_default_prob = bn_default_prob.replace([np.inf, -np.inf], np.nan)
fallback_count = int(bn_default_prob.isna().sum())
bn_default_prob = bn_default_prob.fillna(fallback_probability)

bn_true = (bn_test_data["loan_status"] == "Default").astype(int)
bn_pred = (bn_default_prob >= 0.50).astype(int)

print("Bayesian Network fallback probability:", round(fallback_probability, 4))
print("Bayesian Network fallback predictions used:", fallback_count)

# Report the Bayesian Network as a supporting classification model.
bn_metrics = pd.DataFrame([
    "model_source": "Bayesian Network",
    "model": "pgmpy HillClimbSearch + VariableElimination",
    "roc_auc": roc_auc_score(bn_true, bn_default_prob),
    "accuracy": accuracy_score(bn_true, bn_pred),
    "precision": precision_score(bn_true, bn_pred, zero_division=0),
    "recall": recall_score(bn_true, bn_pred, zero_division=0),
])

```

```

    "f1": f1_score(bn_true, bn_pred, zero_division=0),
  })
bn_metrics.style.format({
    "roc_auc": "{:.4f}", "accuracy": "{:.4f}", "precision": "{:.4f}",
    "recall": "{:.4f}", "f1": "{:.4f}",
  })

```

Bayesian Network fallback probability: 0.2245

Bayesian Network fallback predictions used: 0

[57]: <pandas.io.formats.style.Styler at 0x1a0561be8a0>

The Bayesian Network metrics show how well grouped conditional probabilities classify default on the test data. This model is useful for interpretation, while the Scikit-learn classifiers remain the main predictive benchmark.

1.3 4. Classification Modeling

1.3.1 4.1 Shared Modeling Preprocessing Pipeline

The same preprocessing step is used before both classification and regression. The raw dataframe keeps readable columns such as `income_group`, `loan_intent`, and `person_home_ownership`, but machine learning models cannot directly train on text labels.

The pipeline solves this by converting nominal text columns into one-hot encoded numeric columns. After this conversion, the model receives only numeric model-ready columns. The original text columns stay in the CSV for readability, but they are not passed to the final model matrix.

```

[58]: # Define target names and feature groups used by the models.
classification_target = "loan_status"
loan_amount_target = "loan_amnt"
loan_percent_target = "loan_percent_income"
interest_rate_target = "loan_int_rate"

# Numeric modeling features keep transformed/ratio versions and remove
↳avoidable duplicates.
classification_numeric_cols = [
    "person_age", "person_emp_length", "loan_percent_income", "loan_int_rate",
    "cb_person_cred_hist_length", "log_person_income", "log_loan_amnt",
    "income_to_loan_ratio", "employment_to_age_ratio",
    ↳"credit_history_to_age_ratio",
]
regression_numeric_cols = [
    "person_age", "person_emp_length", "cb_person_cred_hist_length",
    "log_person_income", "employment_to_age_ratio",
    ↳"credit_history_to_age_ratio",
]

```

```

nominal_cols = [
    "person_home_ownership", "loan_intent", "income_group",
    "loan_amount_group", "loan_percent_income_group", "age_group",
]
binary_cols = ["cb_person_default_on_file"]
ordinal_cols = ["loan_grade"]

classification_feature_cols = classification_numeric_cols + nominal_cols +
    ↪ binary_cols + ordinal_cols
regression_feature_cols = regression_numeric_cols + [
    "person_home_ownership", "loan_intent", "income_group",
    "age_group", "cb_person_default_on_file",
]

# Combined numeric list lets the shared preprocessor find numeric columns for
    ↪ either task.
numeric_cols = sorted(set(classification_numeric_cols +
    ↪ regression_numeric_cols))

feature_counts = pd.DataFrame([
    {"feature_set": "classification", "feature_count":
    ↪ len(classification_feature_cols)},
    {"feature_set": "regression", "feature_count":
    ↪ len(regression_feature_cols)},
])
feature_counts

```

```

[58]:      feature_set  feature_count
0  classification         18
1      regression         11

```

```

[59]: def make_one_hot_encoder():
    # Use the correct OneHotEncoder argument name for the installed
    ↪ scikit-learn version.
    try:
        return OneHotEncoder(handle_unknown="ignore", sparse_output=False)
    except TypeError:
        return OneHotEncoder(handle_unknown="ignore", sparse=False)

def make_model_preprocessor(feature_cols):
    # Keep only the columns that are included in the current model feature set.
    model_numeric_cols = [col for col in numeric_cols if col in feature_cols]
    model_nominal_cols = [col for col in nominal_cols if col in feature_cols]
    model_binary_cols = [col for col in binary_cols if col in feature_cols]
    model_ordinal_cols = [col for col in ordinal_cols if col in feature_cols]
    transformers = []

```

```

    # Numeric features are median-imputed and standardized for scale-sensitive
    ↪models.
    if model_numeric_cols:
        transformers.append(("num", Pipeline([
            ("imputer", SimpleImputer(strategy="median")),
            ("scaler", StandardScaler()),
        ]), model_numeric_cols))
    # Nominal categorical features start as readable text columns in df_model.
    # OneHotEncoder replaces each text column with temporary numeric 0/1
    ↪columns.
    # Example: income_group becomes columns like nom__income_group_Low.
    if model_nominal_cols:
        transformers.append(("nom", Pipeline([
            ("imputer", SimpleImputer(strategy="most_frequent")),
            ("onehot", make_one_hot_encoder()),
        ]), model_nominal_cols))
    # The default-on-file flag is encoded with a fixed N/Y order.
    if model_binary_cols:
        transformers.append(("bin", Pipeline([
            ("imputer", SimpleImputer(strategy="most_frequent")),
            ("ordinal", OrdinalEncoder(categories=[["N", "Y"]])),
        ]), model_binary_cols))
    # Loan grade is ordinal because grades A through G represent increasing
    ↪credit risk.
    if model_ordinal_cols:
        transformers.append(("grade", OrdinalEncoder(
            categories=[["A", "B", "C", "D", "E", "F", "G"]],
            handle_unknown="use_encoded_value",
            unknown_value=-1,
        ), model_ordinal_cols))

    # ColumnTransformer combines all preprocessing steps into one reusable
    ↪object.
    return ColumnTransformer(transformers=transformers, remainder="drop")

```

[60]: `DEFAULT_DECISION_THRESHOLD = 0.50`

```

class ThresholdClassifier:
    def __init__(self, model, threshold=DEFAULT_DECISION_THRESHOLD):
        # Store the fitted probability model and the selected decision
        ↪threshold.
        self.model = model
        self.threshold = threshold

    def predict_proba(self, X):

```



```

        # Keep probability output unchanged for ROC-AUC and probability
        ↪reporting.
        return self.model.predict_proba(X)

    def predict(self, X):
        # Convert default probabilities into class labels using the selected
        ↪threshold.
        default_probability = self.predict_proba(X)[: , 1]
        return (default_probability >= self.threshold).astype(int)

def make_pipeline_for(model, feature_cols):
    # Build one reusable pipeline: preprocessing first, model second.
    return Pipeline([
        ("preprocess", make_model_preprocessor(feature_cols)),
        ("model", clone(model)),
    ])

def model_file_stem(model_name):
    # Convert display names into simple file names.
    return model_name.lower().replace(" ", "_").replace("-", "_")

def save_model(model, folder, file_name):
    # Save one fitted model and return its path.
    folder.mkdir(parents=True, exist_ok=True)
    path = folder / file_name
    joblib.dump(model, path)
    return path

def save_model_group(models, folder, prefix=""):
    # Save every fitted model in a dictionary.
    rows = []
    for model_name, fitted_model in models.items():
        file_name = f"{prefix}{model_file_stem(model_name)}.joblib"
        path = save_model(fitted_model, folder, file_name)
        rows.append({"model": model_name, "saved_path": str(path)})
    return pd.DataFrame(rows)

def load_model(folder, file_name):
    # Load one saved model file.
    return joblib.load(folder / file_name)

```

The preprocessing object is shared by both classification and regression. The same transformer pre-

pares numeric and categorical columns, and only the final estimator changes depending on whether the task predicts default status or a numeric lending target. This avoids duplicate preprocessing code and keeps model comparisons consistent.

1.3.2 Encoded Feature Names After One-Hot Encoding

The saved CSV keeps readable category columns such as `income_group`, `loan_intent`, and `person_home_ownership`. During model training, the preprocessing pipeline temporarily changes these text categories into numeric one-hot columns with names that start with `nom__`.

This preview shows what the model receives after preprocessing. It also confirms that the original nominal text columns are not kept in the final encoded model input. No manual deletion from the CSV is needed.

```
[61]: # Fit the classification preprocessor only to inspect the final model-ready
      ↪ columns.
      encoded_feature_preprocessor =
      ↪ make_model_preprocessor(classification_feature_cols)
      encoded_feature_preprocessor.fit(df_model[classification_feature_cols])

      # Get the transformed column names created by scaling, one-hot encoding, and
      ↪ ordinal encoding.
      encoded_feature_names = encoded_feature_preprocessor.get_feature_names_out()
      encoded_feature_names_table = pd.DataFrame({
          "encoded_feature_name": encoded_feature_names
      })

      print("Total encoded model columns:", len(encoded_feature_names))
      encoded_feature_names_table
```

Total encoded model columns: 38

```
[61]:          encoded_feature_name
0          num__cb_person_cred_hist_length
1          num__credit_history_to_age_ratio
2          num__employment_to_age_ratio
3          num__income_to_loan_ratio
4          num__loan_int_rate
5          num__loan_percent_income
6          num__log_loan_amnt
7          num__log_person_income
8          num__person_age
9          num__person_emp_length
10         nom__person_home_ownership_MORTGAGE
11         nom__person_home_ownership_OTHER
12         nom__person_home_ownership_OWN
13         nom__person_home_ownership_RENT
14         nom__loan_intent_DEBTCONSOLIDATION
15         nom__loan_intent_EDUCATION
```

```

16         nom__loan_intent_HOMEIMPROVEMENT
17         nom__loan_intent_MEDICAL
18         nom__loan_intent_PERSONAL
19         nom__loan_intent_VENTURE
20         nom__income_group_High
21         nom__income_group_Low
22         nom__income_group_Lower-Middle
23         nom__income_group_Middle
24         nom__income_group_Upper-Middle
25         nom__loan_amount_group_Large
26         nom__loan_amount_group_Medium
27         nom__loan_amount_group_Small
28         nom__loan_amount_group_Very Large
29         nom__loan_percent_income_group_High Burden
30         nom__loan_percent_income_group_Low Burden
31         nom__loan_percent_income_group_Moderate Burden
32         nom__loan_percent_income_group_Very High Burden
33         nom__age_group_Adult
34         nom__age_group_Experienced Adult
35         nom__age_group_Young
36         bin__cb_person_default_on_file
37         grade__loan_grade

```

The final verification table checks whether the original nominal text columns are still present after preprocessing.

The expected answer is **False** for each original nominal column. This confirms that the model does not train on raw text columns such as `income_group`; it trains on numeric encoded columns created by the preprocessing pipeline.

```

[62]: # Convert encoded feature names to a set for a direct membership check.
encoded_feature_name_set = set(encoded_feature_names)

# Confirm that each original nominal text column is dropped after preprocessing.
# The encoded numeric columns are created inside the pipeline and used by the
↳model.
nominal_encoding_verification = pd.DataFrame({
    "original_nominal_column": nominal_cols,
    "present_as_original_after_preprocessing": [
        col in encoded_feature_name_set for col in nominal_cols
    ],
    "encoded_columns_created": [
        int(encoded_feature_names_table["encoded_feature_name"].str.
↳startswith(f"nom__{col}_").sum())
        for col in nominal_cols
    ],
})

```

```
nominal_encoding_verification
```

```
[62]:      original_nominal_column  present_as_original_after_preprocessing  \
0      person_home_ownership                                           False
1              loan_intent                                           False
2              income_group                                           False
3      loan_amount_group                                           False
4  loan_percent_income_group                                           False
5              age_group                                           False

      encoded_columns_created
0              4
1              6
2              5
3              4
4              4
5              3
```

1.3.3 4.2 ML Classification Models

The classification models estimate the probability that a borrower will default. The main model results use the standard 0.50 cutoff, while Section 4.4 separately tests the 0.24 cutoff to compare business cost and accuracy.

```
[63]: # Split classification features and target into train and test datasets.
X_cls = df_model[classification_feature_cols].copy()
y_cls = df_model[classification_target].copy()
X_train_cls, X_test_cls, y_train_cls, y_test_cls = train_test_split(
    X_cls, y_cls, test_size=TEST_SIZE, stratify=y_cls, random_state=RANDOM_STATE
)
split_summary = pd.DataFrame([
    {"dataset": "train", "rows": X_train_cls.shape[0], "features": X_train_cls.
    ↪shape[1]},
    {"dataset": "test", "rows": X_test_cls.shape[0], "features": X_test_cls.
    ↪shape[1]},
])
split_summary
```

```
[63]:   dataset  rows  features
0   train  23653         18
1    test   5914         18
```

```
[64]: # Define the Scikit-learn classification models for comparison.
classification_models = {
    "Logistic Regression": LogisticRegression(max_iter=1000,
    ↪class_weight="balanced", random_state=RANDOM_STATE),
```

```

    "Random Forest": RandomForestClassifier(n_estimators=300,
↪class_weight="balanced", random_state=RANDOM_STATE, n_jobs=-1),
    "Gradient Boosting": GradientBoostingClassifier(random_state=RANDOM_STATE),
    "CatBoost": CatBoostClassifier(iterations=300, learning_rate=0.05, depth=6,
↪eval_metric="AUC", random_seed=RANDOM_STATE, verbose=0)
}

```

```

[65]: def evaluate_classification_models(models, X_train, X_test, y_train, y_test,
↪feature_cols):
    rows, fitted_models = [], {}
    for name, model in models.items():
        pipe = make_pipeline_for(model, feature_cols)
        pipe.fit(X_train, y_train)
        y_prob = pipe.predict_proba(X_test)[: , 1]
        y_pred = pipe.predict(X_test)
        rows.append({
            "model": name,
            "roc_auc": roc_auc_score(y_test, y_prob),
            "accuracy": accuracy_score(y_test, y_pred),
            "precision": precision_score(y_test, y_pred, zero_division=0),
            "recall": recall_score(y_test, y_pred, zero_division=0),
            "f1": f1_score(y_test, y_pred, zero_division=0),
        })
        fitted_models[name] = pipe
    metrics = pd.DataFrame(rows).sort_values("roc_auc", ascending=False).
↪reset_index(drop=True)
    return metrics, fitted_models

cls_metrics, cls_models = evaluate_classification_models(
    classification_models, X_train_cls, X_test_cls, y_train_cls, y_test_cls,
↪classification_feature_cols
)
best_cls_name = cls_metrics.iloc[0]["model"]
best_cls_model = cls_models[best_cls_name]
y_cls_prob = best_cls_model.predict_proba(X_test_cls)[: , 1]

sklearn_cls_best = cls_metrics.iloc[[0]].copy()
sklearn_cls_best["model_source"] = "Scikit-learn ML Models"

cls_metrics.style.format({
    "roc_auc": "{:.4f}", "accuracy": "{:.4f}", "precision": "{:.4f}",
    "recall": "{:.4f}", "f1": "{:.4f}",
})

```

[65]: <pandas.io.formats.style.Styler at 0x1a05792df40>

All trained classification models are saved in the classification model folder. The best classifier is also saved with a stable file name, `best_default_classifier.joblib`, so the final example prediction can load it easily.

The stable best-classifier file uses the standard 0.50 decision cutoff. The 0.24 threshold is kept as a separate threshold-tuning test because it is a business-cost decision rule, not the main model training rule.

```
[66]: classification_model_files = save_model_group(cls_models,
        ↪CLASSIFICATION_MODEL_DIR)

best_saved_classifier = ThresholdClassifier(
    model=best_cls_model,
    threshold=DEFAULT_DECISION_THRESHOLD,
)

best_classifier_path = save_model(
    best_saved_classifier,
    CLASSIFICATION_MODEL_DIR,
    "best_default_classifier.joblib",
)

classification_model_files = pd.concat([
    classification_model_files,
    pd.DataFrame([{"model": f"Best classifier: {best_cls_name}", "saved_path":
        ↪str(best_classifier_path)}]),
], ignore_index=True)

classification_model_files
```

```
[66]:
```

	model	\		saved_path
0	Logistic Regression			..\data\models\classification\logistic_regress...
1	Random Forest			..\data\models\classification\random_forest.jo...
2	Gradient Boosting			..\data\models\classification\gradient_boostin...
3	CatBoost			..\data\models\classification\catboost.joblib
4	Best classifier: CatBoost			..\data\models\classification\best_default_cla...

The classification table compares ML models using ROC-AUC, accuracy, precision, recall, and F1. The highest ROC-AUC model is selected because it best separates default from non-default cases across probability thresholds.

1.3.4 4.3 Cross-Validation

Cross-validation gives a more stable performance estimate than one train-test split because the model is evaluated across multiple folds. This section uses the normal 0.50 cutoff, matching the main classification model results.

```
[67]: # Run cross-validation to estimate model performance across multiple folds.
cv_model_name = "CatBoost" if "CatBoost" in classification_models else
    ↪ "Gradient Boosting"
cv_base_model = classification_models[cv_model_name]
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=RANDOM_STATE)
cv_pipe = make_pipeline_for(cv_base_model, classification_feature_cols)
auc_scores = cross_val_score(cv_pipe, X_cls, y_cls, cv=cv, scoring="roc_auc")
oof_pred = cross_val_predict(cv_pipe, X_cls, y_cls, cv=cv)
cv_report = classification_report(y_cls, oof_pred, digits=3, output_dict=True)
cv_results = pd.DataFrame([
    "model": cv_model_name,
    "cv_auc_mean": auc_scores.mean(),
    "cv_auc_std": auc_scores.std(),
    "default_recall": cv_report["1"]["recall"],
    "default_precision": cv_report["1"]["precision"],
    "default_f1": cv_report["1"]["f1-score"],
])
cv_results.style.format({
    "cv_auc_mean": "{:.4f}", "cv_auc_std": "{:.4f}",
    "default_recall": "{:.4f}", "default_precision": "{:.4f}", "default_f1": "{:
    ↪ .4f}",
})
```

```
[67]: <pandas.io.formats.style.Styler at 0x1a0490442c0>
```

The cross-validation result checks whether model performance is stable across multiple folds. These values use the standard classification cutoff, while the tuned 0.24 threshold is reviewed separately in the next section.

1.3.5 4.4 Threshold Tuning

The default probability threshold of 0.50 may not be the best lending policy. This step chooses a threshold that minimizes a simple business cost where a missed default is more expensive than a false alarm.

```
[68]: # Evaluate thresholds using a simple lending cost assumption.
COST_FALSE_NEGATIVE = 5
COST_FALSE_POSITIVE = 1
threshold_rows = []
for threshold in np.arange(0.05, 0.96, 0.01):
    y_threshold_pred = (y_cls_prob >= threshold).astype(int)
```

```

false_negative = ((y_threshold_pred == 0) & (y_test_cls == 1)).sum()
false_positive = ((y_threshold_pred == 1) & (y_test_cls == 0)).sum()
total_cost = COST_FALSE_NEGATIVE * false_negative + COST_FALSE_POSITIVE *
↪false_positive
threshold_rows.append({
    "threshold": threshold,
    "false_negative": false_negative,
    "false_positive": false_positive,
    "total_cost": total_cost,
})
threshold_results = pd.DataFrame(threshold_rows)
best_threshold_row = threshold_results.loc[threshold_results["total_cost"].
↪idxmin()]
best_threshold = float(best_threshold_row["threshold"])
best_threshold_row.to_frame("value")

```

```

[68]:
threshold      0.15
false_negative 211.00
false_positive 528.00
total_cost     1583.00

```

The threshold tuning result gives the best cutoff based on the project cost rule, not based only on accuracy. The 0.24 cutoff may reduce accuracy slightly because it flags more borrowers as possible defaults. However, it can reduce business risk because it catches more actual default cases when missed defaults are treated as more expensive than false alarms.

```

[69]: # Use the selected threshold from threshold tuning.
selected_threshold = 0.15
# Apply the tuned threshold only for this business-cost comparison.
y_pred_threshold_024 = (y_cls_prob >= selected_threshold).astype(int)

# Calculate standard classification metrics at the selected threshold.
threshold_024_results = pd.DataFrame([
    "threshold": selected_threshold,
    "accuracy": accuracy_score(y_test_cls, y_pred_threshold_024),
    "precision": precision_score(y_test_cls, y_pred_threshold_024,
↪zero_division=0),
    "recall": recall_score(y_test_cls, y_pred_threshold_024, zero_division=0),
    "f1": f1_score(y_test_cls, y_pred_threshold_024, zero_division=0),
])
threshold_024_results.style.format({
    "threshold": "{:.2f}",
    "accuracy": "{:.4f}",
    "precision": "{:.4f}",
    "recall": "{:.4f}",
    "f1": "{:.4f}",
})

```




```
} )
```

```
[69]: <pandas.io.formats.style.Styler at 0x1a057a25970>
```

The threshold result identifies a default-probability cutoff based on a simple lending cost assumption. This is important because lending decisions may prefer catching more potential defaults even if that creates more false alarms.

```
[70]: # Plot threshold cost to show the selected operating point.
plt.figure(figsize=(9, 5))
plt.plot(threshold_results["threshold"], threshold_results["total_cost"],
         ↪ linewidth=2)
plt.axvline(best_threshold, color="red", linestyle="--", label=f"Best threshold_
         ↪ = {best_threshold:.2f}")
plt.title("Business Cost by Default Threshold")
plt.xlabel("Default Probability Threshold")
plt.ylabel("Total Cost")
plt.legend()
plt.tight_layout()
plt.show()
```



1.4 5. Regression Modeling

1.4.1 5.1 Regression Dataset for Loan Sizing

Regression models are trained only on non-default loans. This creates a more conservative loan sizing benchmark because the models learn from loans that historically performed successfully. The

same regression setup also estimates interest rate as an additional numeric lending benchmark.

```
[71]: # Keep only non-default loans for conservative loan-sizing regression.
df_regression = df_model[df_model[classification_target] == 0].copy().
    ↪reset_index(drop=True)
X_reg = df_regression[regression_feature_cols].copy()
y_loan_amount = df_regression[loan_amount_target].copy()
y_loan_percent = df_regression[loan_percent_target].copy()
y_interest_rate = df_regression[interest_rate_target].copy()
regression_setup_summary = pd.DataFrame([
    "rows": len(df_regression),
    "features": X_reg.shape[1],
    "loan_amount_mean": y_loan_amount.mean(),
    "loan_percent_mean": y_loan_percent.mean(),
    "interest_rate_mean": y_interest_rate.mean(),
])
regression_setup_summary.style.format({
    "loan_amount_mean": "{:.2f}",
    "loan_percent_mean": "{:.4f}",
    "interest_rate_mean": "{:.2f}",
})
```

```
[71]: <pandas.io.formats.style.Styler at 0x1a05799a930>
```

The regression setup summary confirms that loan-sizing models are trained only on non-default loans. This makes the estimates more conservative because the models learn from loans that historically performed successfully.

```
[72]: # Preview regression predictor columns before model training.
X_reg.head()
```

```
[72]:
```

	person_age	person_emp_length	cb_person_cred_hist_length	\
0	21	5.0		2
1	23	2.0		4
2	23	0.0		4
3	25	9.0		2
4	24	0.0		3

	log_person_income	employment_to_age_ratio	credit_history_to_age_ratio	\
0	9.169623	0.238095		0.095238
1	11.652696	0.086957		0.173913
2	11.695255	0.000000		0.173913
3	11.827744	0.360000		0.080000
4	9.303922	0.000000		0.125000

	person_home_ownership	loan_intent	income_group	age_group	\
0	OWN	EDUCATION	Low	Young	
1	RENT	EDUCATION	High	Young	

2	RENT	EDUCATION	High	Young
3	RENT	PERSONAL	High	Young
4	OWN	PERSONAL	Low	Young

	cb_person_default_on_file
0	N
1	N
2	N
3	Y
4	N

```
[73]: # Split regression data into matching train and test datasets.
X_reg_train, X_reg_test, y_amt_train, y_amt_test = train_test_split(
    X_reg, y_loan_amount, test_size=TEST_SIZE, random_state=RANDOM_STATE
)
_, _, y_pct_train, y_pct_test = train_test_split(
    X_reg, y_loan_percent, test_size=TEST_SIZE, random_state=RANDOM_STATE
)
_, _, y_rate_train, y_rate_test = train_test_split(
    X_reg, y_interest_rate, test_size=TEST_SIZE, random_state=RANDOM_STATE
)
```

The raw table may still show words such as RENT or DEBTCONSOLIDATION. This is expected because the pipeline needs the original column names.

The model does not use those text values directly. The preview below shows the encoded regression test data after preprocessing. Category columns become 0/1 columns, while scaled numeric columns can appear as decimal values.

```
[74]: # Fit the shared preprocessor only on regression training features.
regression_preview_preprocessor =
    ↳ make_model_preprocessor(regression_feature_cols)
regression_preview_preprocessor.fit(X_reg_train)

# Transform the regression test features into the numeric matrix used by models.
X_reg_test_encoded = regression_preview_preprocessor.transform(X_reg_test)
if hasattr(X_reg_test_encoded, "toarray"):
    X_reg_test_encoded = X_reg_test_encoded.toarray()

# Use feature names when the installed scikit-learn version supports them.
try:
    regression_encoded_feature_names = regression_preview_preprocessor.
        ↳ get_feature_names_out()
except Exception:
    regression_encoded_feature_names = [f"feature_{i}" for i in
        ↳ range(X_reg_test_encoded.shape[1])]

# Display only a small preview so the table stays readable.
```

```

encoded_preview_column_count = min(20, X_reg_test_encoded.shape[1])
X_reg_test_encoded_preview = pd.DataFrame(
    X_reg_test_encoded[:, :encoded_preview_column_count],
    columns=regression_encoded_feature_names[:encoded_preview_column_count],
)
print("Encoded regression test data shape:", X_reg_test_encoded.shape)
X_reg_test_encoded_preview.head().round(3)

```

Encoded regression test data shape: (4587, 25)

```

[74]:  num__cb_person_cred_hist_length  num__credit_history_to_age_ratio  \
0                                1.413                                2.065
1                               -0.994                               -1.247
2                                0.510                                0.208
3                               -0.994                               -1.174
4                                0.811                                0.998

      num__employment_to_age_ratio  num__log_person_income  num__person_age  \
0                -0.232                -0.065                0.045
1                 0.102                 0.549               -0.177
2                -0.906                 0.127                1.600
3                -0.721                 0.304               -0.621
4               -1.350                -0.789                0.490

      num__person_emp_length  nom__person_home_ownership_MORTGAGE  \
0                -0.205                                0.0
1                 0.058                                0.0
2                -0.732                                0.0
3                -0.732                                1.0
4               -1.258                                1.0

      nom__person_home_ownership_OTHER  nom__person_home_ownership_OWEN  \
0                                0.0                                0.0
1                                0.0                                1.0
2                                0.0                                0.0
3                                0.0                                0.0
4                                0.0                                0.0

      nom__person_home_ownership_RENT  nom__loan_intent_DEBTCONSOLIDATION  \
0                                1.0                                0.0
1                                0.0                                0.0
2                                1.0                                0.0
3                                0.0                                0.0
4                                0.0                                0.0

      nom__loan_intent_EDUCATION  nom__loan_intent_HOMEIMPROVEMENT  \
0                                0.0                                0.0

```

1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	1.0	0.0

	nom__loan_intent_MEDICAL	nom__loan_intent_PERSONAL \
0	0.0	0.0
1	0.0	0.0
2	1.0	0.0
3	0.0	1.0
4	0.0	0.0

	nom__loan_intent_VENTURE	nom__income_group_High	nom__income_group_Low \
0	1.0	0.0	0.0
1	1.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

	nom__income_group_Lower-Middle	nom__income_group_Middle
0	0.0	1.0
1	0.0	1.0
2	0.0	1.0
3	0.0	1.0
4	1.0	0.0

This preview confirms that the regression models receive numbers, not raw text. The one-hot columns are 0/1: 1 means that category is present for that row, and 0 means it is not present. Decimal values are normal for scaled numeric columns such as income, age, and loan amount.

1.4.2 5.2 ML Regression Models

Regression models predict numeric targets. MAE, RMSE, and R-squared are used to judge how well each model predicts loan amount, loan percentage of income, and interest rate.

```
[75]: # Define Scikit-learn regression models, including linear and tree-based
      ↪ approaches.
      regression_models = {
          "Linear Regression": LinearRegression(),
          "Ridge Regression": Ridge(alpha=1.0, random_state=RANDOM_STATE),
          "Random Forest Regressor": RandomForestRegressor(
              n_estimators=300, min_samples_leaf=10,
              random_state=RANDOM_STATE, n_jobs=-1,
          ),
          "Gradient Boosting Regressor": ↪
      ↪ GradientBoostingRegressor(random_state=RANDOM_STATE),
```

```

    "CatBoost Regressor": CatBoostRegressor(iterations=300, learning_rate=0.05,
↳depth=6, loss_function="RMSE", random_seed=RANDOM_STATE, verbose=0)
}

```

```

[76]: def evaluate_regression_models(models, X_train, X_test, y_train, y_test,
↳target_name):
    rows, fitted_models = [], {}
    for name, model in models.items():
        pipe = make_pipeline_for(model, regression_feature_cols)
        pipe.fit(X_train, y_train)
        y_pred = pipe.predict(X_test)
        if target_name == loan_percent_target:
            y_pred = np.clip(y_pred, 0, df_regression[loan_percent_target].
↳quantile(0.99))
        rows.append({
            "target": target_name,
            "model_source": "Scikit-learn ML Models",
            "model": name,
            "mae": mean_absolute_error(y_test, y_pred),
            "rmse": np.sqrt(mean_squared_error(y_test, y_pred)),
            "r2": r2_score(y_test, y_pred),
        })
        fitted_models[name] = pipe
    metrics = pd.DataFrame(rows).sort_values("rmse").reset_index(drop=True)
    return metrics, fitted_models

```

```

[77]: regression_targets = {
    loan_amount_target: y_amt_train,
    loan_percent_target: y_pct_train,
    interest_rate_target: y_rate_train,
}
regression_tests = {
    loan_amount_target: y_amt_test,
    loan_percent_target: y_pct_test,
    interest_rate_target: y_rate_test,
}

regression_metrics_by_target = {}
regression_models_by_target = {}
for target_name, y_train_target in regression_targets.items():
    metrics, fitted_models = evaluate_regression_models(
        regression_models,
        X_reg_train,
        X_reg_test,
        y_train_target,
        regression_tests[target_name],
        target_name,
    )

```

```

    )
    regression_metrics_by_target[target_name] = metrics
    regression_models_by_target[target_name] = fitted_models

loan_amount_metrics = regression_metrics_by_target[loan_amount_target]
loan_percent_metrics = regression_metrics_by_target[loan_percent_target]
interest_rate_metrics = regression_metrics_by_target[interest_rate_target]

loan_amount_models = regression_models_by_target[loan_amount_target]
loan_percent_models = regression_models_by_target[loan_percent_target]
interest_rate_models = regression_models_by_target[interest_rate_target]

```

```
[78]: loan_amount_metrics.style.format({"mae": "{:.2f}", "rmse": "{:.2f}", "r2": "{:.4f}"})
```

```
[78]: <pandas.io.formats.style.Styler at 0x1a0579fa240>
```

```
[79]: loan_percent_metrics.style.format({"mae": "{:.4f}", "rmse": "{:.4f}", "r2": "{:.4f}"})
```

```
[79]: <pandas.io.formats.style.Styler at 0x1a055cbefc0>
```

```
[80]: interest_rate_metrics.style.format({"mae": "{:.4f}", "rmse": "{:.4f}", "r2": "{:.4f}"})
```

```
[80]: <pandas.io.formats.style.Styler at 0x1a057a0cda0>
```

1.4.3 5.3 Regression Metric Interpretation

MAE is the average absolute prediction error. Lower MAE means the model is closer to the actual value on average.

RMSE is similar to MAE, but it penalizes large mistakes more strongly. Lower RMSE is better.

R² shows how much variation in the target is explained by the model. Higher R² is better, and values closer to 1 indicate stronger predictive performance.

```
[81]: def r2_strength(r2_value):
    # Convert R-squared into a simple project-level interpretation.
    if r2_value >= 0.70:
        return "strong"
    if r2_value >= 0.40:
        return "moderate"
    return "weak"

def regression_conclusion(target_name, best_model, r2_value):
```

```

# Make weak regression performance clear in the conclusion table.
strength = r2_strength(r2_value)
if strength == "weak":
    return (
        f"{target_name}: {best_model} has weak prediction performance based
on R-squared. "
        "The estimate should be treated as a rough benchmark, not a strong
prediction."
    )
    return (
        f"{target_name}: {best_model} has {strength} prediction performance
based on R-squared."
    )

best_amount_row = loan_amount_metrics.iloc[0]
best_percent_row = loan_percent_metrics.iloc[0]
best_rate_row = interest_rate_metrics.iloc[0]
amount_mae_ratio = best_amount_row["mae"] / y_amt_test.mean()
amount_rmse_ratio = best_amount_row["rmse"] / y_amt_test.mean()
percent_mae_ratio = best_percent_row["mae"] / y_pct_test.mean()
percent_rmse_ratio = best_percent_row["rmse"] / y_pct_test.mean()
rate_mae_ratio = best_rate_row["mae"] / y_rate_test.mean()
rate_rmse_ratio = best_rate_row["rmse"] / y_rate_test.mean()

regression_interpretation = pd.DataFrame([
    {
        "target": loan_amount_target,
        "best_model": best_amount_row["model"],
        "mae": best_amount_row["mae"],
        "rmse": best_amount_row["rmse"],
        "r2": best_amount_row["r2"],
        "mae_vs_target_mean": amount_mae_ratio,
        "rmse_vs_target_mean": amount_rmse_ratio,
        "performance_summary": f"{r2_strength(best_amount_row['r2']).title()}
explanatory performance by R-squared.",
        "conclusion": regression_conclusion(loan_amount_target,
best_amount_row["model"], best_amount_row["r2"]),
    },
    {
        "target": loan_percent_target,
        "best_model": best_percent_row["model"],
        "mae": best_percent_row["mae"],
        "rmse": best_percent_row["rmse"],
        "r2": best_percent_row["r2"],
        "mae_vs_target_mean": percent_mae_ratio,
        "rmse_vs_target_mean": percent_rmse_ratio,
    }
])

```



```

        "performance_summary": f"{r2_strength(best_percent_row['r2']).title()}\n
↳explanatory performance by R-squared.",
        "conclusion": regression_conclusion(loan_percent_target,\n
↳best_percent_row["model"], best_percent_row["r2"]),
    },
    {
        "target": interest_rate_target,
        "best_model": best_rate_row["model"],
        "mae": best_rate_row["mae"],
        "rmse": best_rate_row["rmse"],
        "r2": best_rate_row["r2"],
        "mae_vs_target_mean": rate_mae_ratio,
        "rmse_vs_target_mean": rate_rmse_ratio,
        "performance_summary": f"{r2_strength(best_rate_row['r2']).title()}\n
↳explanatory performance by R-squared.",
        "conclusion": regression_conclusion(interest_rate_target,\n
↳best_rate_row["model"], best_rate_row["r2"]),
    }
])

# Display the metric interpretation and final regression conclusion for each\n
↳target.
regression_interpretation.style.format({
    "mae": "{:.4f}", "rmse": "{:.4f}", "r2": "{:.4f}",
    "mae_vs_target_mean": "{:.2%}", "rmse_vs_target_mean": "{:.2%}",
})

```

[81]: <pandas.io.formats.style.Styler at 0x1a057306180>

The best regression models are selected here for the later applicant-level loan estimate. This keeps model selection separate from plots and makes the inference example easier to follow.

```

[82]: best_regression = {
    loan_amount_target: loan_amount_metrics.iloc[0]["model"],
    loan_percent_target: loan_percent_metrics.iloc[0]["model"],
    interest_rate_target: interest_rate_metrics.iloc[0]["model"],
}

best_amount_name = best_regression[loan_amount_target]
best_percent_name = best_regression[loan_percent_target]
best_rate_name = best_regression[interest_rate_target]
best_amount_model = loan_amount_models[best_amount_name]
best_percent_model = loan_percent_models[best_percent_name]
best_rate_model = interest_rate_models[best_rate_name]

selected_regression_models = pd.DataFrame([
    {"target": target, "selected_model": model_name}

```

```

        for target, model_name in best_regression.items()
    ])
    selected_regression_models

```

```

[82]:
      target      selected_model
0      loan_amnt  Gradient Boosting Regressor
1  loan_percent_income  Gradient Boosting Regressor
2      loan_int_rate      CatBoost Regressor

```

All trained regression models are saved in the regression model folder. The best model for each target is also saved with a stable file name so the final applicant loan estimate can load those best models easily.

```

[83]: regression_model_files = []
for target_name, fitted_models in regression_models_by_target.items():
    saved = save_model_group(fitted_models, REGRESSION_MODEL_DIR,
        prefix=f"{target_name}_")
    saved.insert(0, "target", target_name)
    regression_model_files.append(saved)
regression_model_files = pd.concat(regression_model_files, ignore_index=True)

best_regression_files = [
    (loan_amount_target, best_amount_name, best_amount_model,
    ↪ "best_loan_amount_regressor.joblib"),
    (loan_percent_target, best_percent_name, best_percent_model,
    ↪ "best_loan_percent_income_regressor.joblib"),
    (interest_rate_target, best_rate_name, best_rate_model,
    ↪ "best_interest_rate_regressor.joblib"),
]
for target_name, model_name, fitted_model, file_name in best_regression_files:
    path = save_model(fitted_model, REGRESSION_MODEL_DIR, file_name)
    regression_model_files = pd.concat([
        regression_model_files,
        pd.DataFrame([{"target": target_name, "model": f"Best regressor:
    ↪ {model_name}", "saved_path": str(path)}]),
    ], ignore_index=True)

regression_model_files

```

```

[83]:
      target      model \
0      loan_amnt      Linear Regression
1      loan_amnt      Ridge Regression
2      loan_amnt      Random Forest Regressor
3      loan_amnt      Gradient Boosting Regressor
4      loan_amnt      CatBoost Regressor
5  loan_percent_income      Linear Regression
6  loan_percent_income      Ridge Regression

```

```

7   loan_percent_income      Random Forest Regressor
8   loan_percent_income      Gradient Boosting Regressor
9   loan_percent_income      CatBoost Regressor
10  loan_int_rate             Linear Regression
11  loan_int_rate             Ridge Regression
12  loan_int_rate             Random Forest Regressor
13  loan_int_rate             Gradient Boosting Regressor
14  loan_int_rate             CatBoost Regressor
15  loan_amnt                 Best regressor: Gradient Boosting Regressor
16  loan_percent_income       Best regressor: Gradient Boosting Regressor
17  loan_int_rate             Best regressor: CatBoost Regressor

```

```

                                saved_path
0   ..\data\models\regression\loan_amnt_linear_reg...
1   ..\data\models\regression\loan_amnt_ridge_regr...
2   ..\data\models\regression\loan_amnt_random_for...
3   ..\data\models\regression\loan_amnt_gradient_b...
4   ..\data\models\regression\loan_amnt_catboost_r...
5   ..\data\models\regression\loan_percent_income_...
6   ..\data\models\regression\loan_percent_income_...
7   ..\data\models\regression\loan_percent_income_...
8   ..\data\models\regression\loan_percent_income_...
9   ..\data\models\regression\loan_percent_income_...
10  ..\data\models\regression\loan_int_rate_linear...
11  ..\data\models\regression\loan_int_rate_ridge_...
12  ..\data\models\regression\loan_int_rate_random...
13  ..\data\models\regression\loan_int_rate_gradie...
14  ..\data\models\regression\loan_int_rate_catboo...
15  ..\data\models\regression\best_loan_amount_reg...
16  ..\data\models\regression\best_loan_percent_in...
17  ..\data\models\regression\best_interest_rate_r...

```

1.5 6. Model Comparison and Applicant Inference

1.5.1 6.1 Model Comparison

The tables below summarize the final classification and regression results used in this project.

1.5.2 6.1.1 Classification Model Comparison

```

[84]: classification_compare_parts = [sklearn_cls_best[["model_source", "model", "
    ↪"roc_auc", "accuracy", "precision", "recall", "f1"]]]
if bn_metrics is not None:
    classification_compare_parts.append(bn_metrics[["model_source", "model", "
    ↪"roc_auc", "accuracy", "precision", "recall", "f1"]])

```

```

classification_comparison = pd.concat(classification_compare_parts,
    ↪ ignore_index=True)

# Display Scikit-learn and Bayesian Network classification results together.
classification_comparison.style.format({
    "roc_auc": "{:.4f}", "accuracy": "{:.4f}", "precision": "{:.4f}",
    "recall": "{:.4f}", "f1": "{:.4f}",
})

```

[84]: <pandas.io.formats.style.Styler at 0x1a05722f2c0>

1.5.3 6.1.2 Regression Model Comparison: Loan Amount

```

[85]: amount_model_comparison = loan_amount_metrics.iloc[[0]][["target",
    ↪ "model_source", "model", "mae", "rmse", "r2"]].copy()

# Display the best Scikit-learn model for loan amount.
amount_model_comparison.style.format({"mae": "{:.2f}", "rmse": "{:.2f}", "r2":
    ↪ "{:.4f}"})

```

[85]: <pandas.io.formats.style.Styler at 0x1a056186450>

1.5.4 6.1.3 Regression Model Comparison: Loan Percent Income

```

[86]: percent_model_comparison = loan_percent_metrics.iloc[[0]][["target",
    ↪ "model_source", "model", "mae", "rmse", "r2"]].copy()

# Display the best Scikit-learn model for loan percent income.
percent_model_comparison.style.format({"mae": "{:.4f}", "rmse": "{:.4f}", "r2":
    ↪ "{:.4f}"})

```

[86]: <pandas.io.formats.style.Styler at 0x1a05722fa70>

1.5.5 6.1.4 Regression Model Comparison: Interest Rate

```

[87]: interest_rate_model_comparison = interest_rate_metrics.iloc[[0]][["target",
    ↪ "model_source", "model", "mae", "rmse", "r2"]].copy()

# Display the best Scikit-learn model for interest rate.
interest_rate_model_comparison.style.format({"mae": "{:.4f}", "rmse": "{:.4f}",
    ↪ "r2": "{:.4f}"})

```

[87]: <pandas.io.formats.style.Styler at 0x1a05722f2c0>

1.5.6 6.2 Example Applicant Default Prediction

This example loads the saved best classification model and the saved Bayesian Network model, then applies both to one applicant record. The Scikit-learn model uses the same engineered features used during training. The Bayesian Network uses grouped categorical evidence because its probability tables are learned from discrete variables.

The saved Scikit-learn classifier uses the standard 0.50 cutoff for the final applicant prediction. The tuned 0.24 threshold was reviewed separately as a business-cost comparison in Section 4.4.

```
[88]: example_applicant = {
    "person_age": 28,
    "person_income": 65000,
    "person_home_ownership": "RENT",
    "person_emp_length": 5,
    "loan_intent": "DEBTCONSOLIDATION",
    "loan_grade": "G",
    "loan_amnt": 12000,
    "loan_int_rate": 11.0,
    "loan_percent_income": 12000 / 65000,
    "loan_status": 0,
    "cb_person_default_on_file": "N",
    "cb_person_cred_hist_length": 6,
}

def prepare_applicant_for_classification(applicant_dict):
    # Convert one applicant dictionary into a dataframe.
    applicant_df = pd.DataFrame([applicant_dict])
    # Apply the same feature engineering used before model training.
    applicant_df = add_credit_features(applicant_df)
    # Keep only the predictor columns expected by the classification pipeline.
    return applicant_df[classification_feature_cols]

def prepare_applicant_for_bayesian_network(applicant_dict):
    # Convert the applicant into the grouped/discrete format used by the
    Bayesian Network.
    applicant_df = pd.DataFrame([applicant_dict])
    applicant_df = add_credit_features(applicant_df)
    bn_applicant = applicant_df[bn_cols].copy()
    # Match the target labels used when the Bayesian Network was trained.
    bn_applicant["loan_status"] = bn_applicant["loan_status"].map({0:
    "Non-Default", 1: "Default"})
    return bn_applicant.astype(str).iloc[0]

def default_label(probability, threshold=0.50):
```

```

    # Convert a probability into the same decision label used in the
    ↪classification task.
    return "Default" if probability >= threshold else "Non-Default"

```

```

[89]: # Load the saved best classifier and saved Bayesian Network model for
    ↪prediction.
saved_best_cls_model = load_model(CLASSIFICATION_MODEL_DIR,
    ↪"best_default_classifier.joblib")
saved_bn_model = load_model(BAYESIAN_MODEL_DIR, "bayesian_network_model.joblib")
saved_bn_metadata = joblib.load(BAYESIAN_MODEL_DIR / "bayesian_network_metadata.
    ↪joblib")
applicant_cls_features = prepare_applicant_for_classification(example_applicant)

# Predict default probability and class using the saved 0.24 threshold-aware
    ↪classifier.
sklearn_default_probability = float(saved_best_cls_model.
    ↪predict_proba(applicant_cls_features)[: , 1][0])
sklearn_predicted_class = int(saved_best_cls_model.
    ↪predict(applicant_cls_features)[0])

# Prepare grouped evidence for Bayesian Network inference.
bn_applicant_row = prepare_applicant_for_bayesian_network(example_applicant)
bayesian_default_probability = predict_default_probability_with_bn(
    saved_bn_model,
    bn_applicant_row,
    target_col=saved_bn_metadata["target_col"],
    fallback_probability=saved_bn_metadata["fallback_probability"],
)
bayesian_predicted_class = int(bayesian_default_probability >= 0.50)

# Present both applicant-level default predictions in one compact table.
applicant_default_prediction = pd.DataFrame([
    {
        "model_source": "Scikit-learn ML",
        "model": best_cls_name,
        "default_probability": sklearn_default_probability,
        "predicted_class": sklearn_predicted_class,
        "decision_threshold": saved_best_cls_model.threshold,
        "predicted_label": "Default" if sklearn_predicted_class == 1 else
    ↪"Non-Default",
    },
    {
        "model_source": "Bayesian Network",
        "model": "pgmpy HillClimbSearch + VariableElimination",
        "default_probability": bayesian_default_probability,
        "predicted_class": bayesian_predicted_class,

```

```

        "decision_threshold": 0.50,
        "predicted_label": default_label(bayesian_default_probability),
    },
])
applicant_default_prediction.style.format({
    "default_probability": "{:.2%}",
    "decision_threshold": "{:.2f}",
})

```

[89]: <pandas.io.formats.style.Styler at 0x1a05795f590>

The applicant default-prediction table gives two probability-based views for the same applicant. The Scikit-learn ML row uses the saved best classifier with the standard 0.50 cutoff. The Bayesian Network row uses its probability estimate as a separate supporting comparison.

1.5.7 6.3 Example Applicant Loan Estimate

This example loads the saved best regression model for each target and estimates loan amount, loan percentage of income, and interest rate for one applicant. The conservative recommendation is the lower value between the direct loan amount estimate and the income-based estimate.

All regression models were saved earlier, but this final example uses the stable best-model files for the final estimate.

```

[90]: def prepare_applicant_for_regression(applicant_dict):
    # Convert one applicant dictionary into a dataframe.
    applicant_df = pd.DataFrame([applicant_dict])
    # Apply the same feature engineering used during model training.
    applicant_df = add_credit_features(applicant_df)
    # Keep only the regression predictor columns expected by the trained
    ↪ pipelines.
    return applicant_df[regression_feature_cols]

def estimate_loan_for_applicant(applicant_dict, amount_model, percent_model,
    ↪ rate_model):
    # Prepare the applicant record in the same format as the regression
    ↪ training data.
    applicant_features = prepare_applicant_for_regression(applicant_dict)
    # Predict loan amount directly from the loan amount regression model.
    predicted_amount = float(amount_model.predict(applicant_features)[0])

    # Predict loan percent of income from the loan burden regression model.
    predicted_percent = float(percent_model.predict(applicant_features)[0])

    max_reasonable_percent = float(df_regression[loan_percent_target].
    ↪ quantile(0.99))

```

```

    # Clip extreme percentage predictions to the historical non-default range.
    predicted_percent = float(np.clip(predicted_percent, 0,
↪max_reasonable_percent))

    # Predict interest rate as an informational estimate.
    predicted_interest_rate = float(rate_model.predict(applicant_features)[0])

    # Convert the predicted percentage into an income-based dollar amount.
    income_based_amount = float(applicant_dict["person_income"]) *
↪predicted_percent
    # Use the smaller amount as the conservative recommendation.
    conservative_amount = min(predicted_amount, income_based_amount)
    return {
        "predicted_loan_amount": round(predicted_amount, 2),
        "predicted_loan_percent_income": round(predicted_percent, 4),
        "predicted_interest_rate": round(predicted_interest_rate, 4),
        "income_based_amount": round(income_based_amount, 2),
        "conservative_recommended_amount": round(conservative_amount, 2),
    }

```

```

[91]: # Load the saved best regression models for applicant-level prediction.
saved_amount_model = load_model(REGRESSION_MODEL_DIR,
↪"best_loan_amount_regressor.joblib")
saved_percent_model = load_model(REGRESSION_MODEL_DIR,
↪"best_loan_percent_income_regressor.joblib")
saved_rate_model = load_model(REGRESSION_MODEL_DIR,
↪"best_interest_rate_regressor.joblib")

# Reuse the same example applicant from the default-prediction section.
loan_estimate = estimate_loan_for_applicant(example_applicant,
↪saved_amount_model, saved_percent_model, saved_rate_model)
loan_estimate_table = pd.DataFrame([loan_estimate])

# Show the final applicant estimate table.
loan_estimate_table

```

```

[91]: predicted_loan_amount predicted_loan_percent_income \
0                8943.68                0.1377

    predicted_interest_rate income_based_amount \
0                10.1349                8952.72

    conservative_recommended_amount
0                8943.68

```

The loan-estimate table combines direct loan amount prediction with an income-based estimate. The conservative recommendation uses the lower of those two values so the final amount is not

higher than the model's affordability estimate.